

modelling_clustering_v2

August 12, 2025

0.0.1 Clustering

```
[2]: # ensure scikit-learn's runs single-threaded to reduce memory load.
os.environ["OMP_NUM_THREADS"] = "1" # OpenMP threads
os.environ["MKL_NUM_THREADS"] = "1" # MKL's own thread pool
```

```
[ ]: from dotenv import load_dotenv
import os
```

Load data

```
[ ]: load_dotenv(".env") # defaults to loading a `.env` file in the current_
    ↪directory
```

```
host = os.getenv("MY_SQL_HOST")
user = os.getenv("MY_SQL_USER")
password = os.getenv("MY_SQL_PASSWORD")
database = os.getenv("MY_SQL_DATABASE")
```

```
[ ]: import mysql.connector

def get_connection(host, user, password, database):
    mydb = mysql.connector.connect(
        host=host,
        user=user,
        password=password,
        database=database
    )
    return mydb
```

```
[ ]: connection = get_connection(host=host, user=user, password=password,
    ↪database=database)
cursor = connection.cursor()
root_location = "E:\\applied data science_
    ↪capstone\\anime-recommendation\\modelling\\clustering\\exploration"
```

```
[ ]: query = """
SELECT
    user_id,
```

```

        anime_id,
        rating
FROM
    rating_v2
ORDER BY
    user_id;
"""
cursor.execute(query)
raw_data = cursor.fetchall()

```

```

[ ]: import pandas as pd

df = pd.DataFrame(raw_data, columns=["user_id", "anime_id", "rating"])
df["rating"] = pd.to_numeric(df["rating"])
# save data to file
filename = "E:\\applied data science\\
↳capstone\\data\\etl\\extract\\transactions\\clustering_13_Jul.csv"
df.to_csv(filename, index=False)

```

```

[5]: df = pd.read_csv(filename)
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 47161882 entries, 0 to 47161881
Data columns (total 3 columns):
#   Column      Dtype
---  -
0   user_id     int64
1   anime_id    int64
2   rating      float64
dtypes: float64(1), int64(2)
memory usage: 1.1 GB

```

```

[12]: df.isna().sum()

```

```

[12]: user_id      0
anime_id      0
rating        0
dtype: int64

```

```

[7]: from sklearn.model_selection import train_test_split

```

```

[ ]: # split into train, test and validate
train_df, test_validate_df = train_test_split(df, test_size=0.25,
↳random_state=42)
test_df, validate_df = train_test_split(test_validate_df, test_size=0.2,
↳random_state=42)

```

```
[9]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 35371411 entries, 27745566 to 21081788
Data columns (total 3 columns):
#   Column      Dtype
---  ---
0   user_id     int64
1   anime_id    int64
2   rating      float64
dtypes: float64(1), int64(2)
memory usage: 1.1 GB
```

```
[10]: test_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 9432376 entries, 36901731 to 21750476
Data columns (total 3 columns):
#   Column      Dtype
---  ---
0   user_id     int64
1   anime_id    int64
2   rating      float64
dtypes: float64(1), int64(2)
memory usage: 287.9 MB
```

```
[11]: validate_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 2358095 entries, 39414737 to 31165642
Data columns (total 3 columns):
#   Column      Dtype
---  ---
0   user_id     int64
1   anime_id    int64
2   rating      float64
dtypes: float64(1), int64(2)
memory usage: 72.0 MB
```

```
[13]: print(f"train percentage: {train_df.shape[0]/df.shape[0]}")
      print(f"test percentage: {test_df.shape[0]/df.shape[0]}")
      print(f"validate percentage: {validate_df.shape[0]/df.shape[0]}")
```

```
train percentage: 0.7499999893982178
test percentage: 0.19999999151857425
validate percentage: 0.05000001908320792
```

```
[14]: # save validate df for future user
```

```
filename = "E:\\applied data science\\
↳capstone\\data\\etl\\extract\\transactions\\validate_8_Aug.csv"
validate_df.to_csv(filename, index=False)
```

```
[24]: anime_cat = train_df['anime_id'].astype('category') # create category
anime_codes = anime_cat.cat.codes # get code for categories
index_to_anime = dict(enumerate(anime_cat.cat.categories)) # create mapping for
↳index and anime
anime_to_index = {v: k for k, v in index_to_anime.items()} # reverse mapping
↳for index and anime
```

```
[25]: user_cat = train_df['user_id'].astype('category') # create category
user_codes = user_cat.cat.codes # get code for categories
index_to_user = dict(enumerate(user_cat.cat.categories)) # create mapping for
↳index and user
user_to_index = {v: k for k, v in index_to_user.items()} # reverse mapping for
↳index and user
```

```
[ ]: from scipy.sparse import coo_matrix
# build coo (non-zero values along with their coordinates), then convert to csr
↳(format for sparse matrix)
num_users = user_codes.max() + 1
num_animes = anime_codes.max() + 1
data = train_df['rating'].values
user_item_sparse = coo_matrix((data, (user_codes, anime_codes)),
↳shape=(num_users, num_animes)).tocsr()
```

Training K = 2

```
[27]: from sklearn.cluster import KMeans
```

```
[223]: # evaluation metrics
inertia = []
silhouette_scores = []
davies_bouldin_scores = []
precision_scores = []
recall_scores = []
```

```
[224]: # create cluster
k = 2
kmeans = KMeans(n_clusters=k, random_state=42, max_iter=300)
cluster_labels = kmeans.fit_predict(user_item_sparse)

# inertia measures how internally coherent the k-means clusters are
# lower inertia means points are closer to the centroid
# higher inertia means points are further from the centroid
inertia.append(kmeans.inertia_)
```

Assessing $K = 2$

```
[225]: import numpy as np

[226]: # extract sample
random_generator = np.random.default_rng(42)
sample_indices = random_generator.choice(user_item_sparse.shape[0], size=3000,
    ↪replace=False)
sample_matrix = user_item_sparse[sample_indices]
sample_labels = cluster_labels[sample_indices]
sample_index_to_user = { k:v for k, v in index_to_user.items() if k in
    ↪sample_indices}

[227]: from sklearn.metrics import silhouette_score
from sklearn.metrics import davies_bouldin_score
```

Silhouette

```
[228]: # silhouette score is used to determine the effectiveness of the clusters
# close to -1 means points in cluster are closer to points in other clusters
# close to 1 means the points are close to other points in the same cluster
sil_score = silhouette_score(sample_matrix, sample_labels, metric='cosine')
print(f"Silhouette Score: {sil_score:.3f}")
silhouette_scores.append(sil_score)
```

Silhouette Score: -0.002

Davies-Bouldin Index

```
[229]: # davies-bouldin index is used to compare how well separated a cluster is to
    ↪its size
# lower score is better as it means the clusters are more compact and the
    ↪points are far apart from centroids
db_score = davies_bouldin_score(sample_matrix.toarray(), sample_labels)
print(f"Davies-Bouldin Index: {db_score:.3f}")
davies_bouldin_scores.append(db_score)
```

Davies-Bouldin Index: 3.621

Visualize clusters

```
[230]: from sklearn.decomposition import PCA
import matplotlib.pyplot as plt
import seaborn as sns
plots_location = "E:\\applied data science\\
    ↪capstone\\anime-recommendation\\modelling\\clustering\\assessing"

[231]: # visualize inspect the clusters
pca = PCA(n_components=2)
user_2d = pca.fit_transform(sample_matrix.toarray())


df_plot = pd.DataFrame({
```

```

    'x': user_2d[:, 0],
    'y': user_2d[:, 1],
    'cluster': sample_labels
})

plt.figure(figsize=(10, 6))
sns.scatterplot(data=df_plot, x='x', y='y', hue='cluster', palette='tab10', s=50)
plt.title("User Clusters Visualized via PCA")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.legend(title='Cluster')
plt.savefig(f"{plots_location}\\plots\\user_clusters_k_2.png")
plt.show()

```



Precision@2 and Recall@2

```

[232]: # assign cluster to users
clusters_df = pd.DataFrame({
    'user_id' : [index_to_user[i] for i in range(len(cluster_labels))],
    'cluster': cluster_labels
})

clusters_train_df = train_df.merge(clusters_df, on="user_id")

```

```
[233]: clusters_train_df.head()
```

```
[233]:
```

	user_id	anime_id	rating	cluster
0	287522	522	8.0	1
1	447249	2033	7.0	1
2	446384	31318	10.0	0
3	400930	3712	8.0	0
4	153921	37520	8.0	0

```
[234]: # calculate mean rating for each anime across a cluster
cluster_mean_ratings_df = clusters_train_df.drop(columns=["user_id"]) \
    .groupby(["cluster", "anime_id"]).mean().reset_index()
```

```
[235]: def precision_recall_at_k(user_id, k=5):
    # retrieve cluster for user
    cluster_id = clusters_train_df.loc[clusters_train_df["user_id"] == user_id,
    ↪'cluster'].values[0]
    cluster_profile = cluster_mean_ratings_df.
    ↪loc[cluster_mean_ratings_df["cluster"] == cluster_id] \
        .drop('cluster', errors='ignore')

    # sort by rating
    recommendations = cluster_profile.sort_values(by=["rating"],
    ↪ascending=False)

    # anime rated by user in test df and train df
    user_test_anime = test_df.loc[test_df["user_id"] == user_id, "anime_id"].
    ↪tolist()
    user_train_anime = train_df[train_df["user_id"] == user_id]

    # recommended top-k
    recommended = recommendations.loc[~recommendations["anime_id"] \
        .isin(user_train_anime["anime_id"].unique()), "anime_id"].tolist()
    recommended_k = recommended[:k]

    # no test data for user
    if len(user_test_anime) == 0:
        return None, None

    precision = len(set(recommended_k) & set(user_test_anime)) / k
    recall = len(set(recommended_k) & set(user_test_anime)) /
    ↪len(user_test_anime)
    return precision, recall
```

```
[236]: precisions, recalls = [], []
for user_id in sample_index_to_user.values():
    p, r = precision_recall_at_k(user_id, k=5)
```

```

    if p is not None:
        precisions.append(p)
        recalls.append(r)

mean_precision = np.mean(precisions)
mean_recall = np.mean(recalls)

precision_scores.append(mean_precision)
recall_scores.append(mean_recall)

print(f"Precision@2: {mean_precision:.6f}")
print(f"Recall@2: {mean_recall:.6f}")

```

Precision@2: 0.017795
Recall@2: 0.001476

Training K = 3

```

[237]: # create cluster
k = 3
kmeans = KMeans(n_clusters=k, random_state=42, max_iter=300)
cluster_labels = kmeans.fit_predict(user_item_sparse)

inertia.append(kmeans.inertia_)

```

Assessing K = 3

```

[238]: sample_labels = cluster_labels[sample_indices]

```

Silhouette

```

[239]: sil_score = silhouette_score(sample_matrix, sample_labels, metric='cosine')
print(f"Silhouette Score: {sil_score:.3f}")
silhoutte_scores.append(sil_score)

```

Silhouette Score: -0.023

Davies-Bouldin Index

```

[240]: db_score = davies_bouldin_score(sample_matrix.toarray(), sample_labels)
print(f"Davies-Bouldin Index: {db_score:.3f}")
davies_bouldin_scores.append(db_score)

```

Davies-Bouldin Index: 4.691

Visualize clusters

```

[241]: # visualize inspect the clusters
pca = PCA(n_components=2)
user_2d = pca.fit_transform(sample_matrix.toarray())

df_plot = pd.DataFrame({
    'x': user_2d[:, 0],

```




```

        'y': user_2d[:, 1],
        'cluster': sample_labels
    })

plt.figure(figsize=(10, 6))
sns.scatterplot(data=df_plot, x='x', y='y', hue='cluster', palette='tab10',
               s=50)
plt.title("User Clusters Visualized via PCA")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.legend(title='Cluster')
plt.savefig(f"{plots_location}\\plots\\user_clusters_k_{k}.png")
plt.show()

```



Precision@3 and Recall@3

```

[242]: # assign cluster to users
clusters_df = pd.DataFrame({
    'user_id' : [index_to_user[i] for i in range(len(cluster_labels))],
    'cluster': cluster_labels
})

clusters_train_df = train_df.merge(clusters_df, on="user_id")

```

```

[243]: clusters_train_df.head()

```

```
[243]:
```

	user_id	anime_id	rating	cluster
0	287522	522	8.0	2
1	447249	2033	7.0	1
2	446384	31318	10.0	2
3	400930	3712	8.0	0
4	153921	37520	8.0	0

```
[244]: # calculate mean rating for each anime across a cluster
cluster_mean_ratings_df = clusters_train_df.drop(columns=["user_id"]) \
    .groupby(["cluster", "anime_id"]).mean().reset_index()
```

```
[245]: precisions, recalls = [], []
for user_id in sample_index_to_user.values():
    p, r = precision_recall_at_k(user_id, k=5)
    if p is not None:
        precisions.append(p)
        recalls.append(r)

mean_precision = np.mean(precisions)
mean_recall = np.mean(recalls)

precision_scores.append(mean_precision)
recall_scores.append(mean_recall)

print(f"Precision@3: {mean_precision:.6f}")
print(f"Recall@3: {mean_recall:.6f}")
```

Precision@3: 0.007657

Recall@3: 0.000430

Training K = 4

```
[246]: # create cluster
k = 4
kmeans = KMeans(n_clusters=k, random_state=42, max_iter=300)
cluster_labels = kmeans.fit_predict(user_item_sparse)

inertia.append(kmeans.inertia_)
```

Assessing K = 4

```
[247]: sample_labels = cluster_labels[sample_indices]
```

Silhouette

```
[248]: sil_score = silhouette_score(sample_matrix, sample_labels, metric='cosine')
print(f"Silhouette Score: {sil_score:.6f}")
silhoutte_scores.append(sil_score)
```

Silhouette Score: -0.024993

Davies-Bouldin Index

```
[249]: db_score = davies_bouldin_score(sample_matrix.toarray(), sample_labels)
print(f"Davies-Bouldin Index: {db_score:.6f}")
davies_bouldin_scores.append(db_score)
```


Davies-Bouldin Index: 4.701756

Visualize clusters

```
[250]: # visualize inspect the clusters
pca = PCA(n_components=2)
user_2d = pca.fit_transform(sample_matrix.toarray())

df_plot = pd.DataFrame({
    'x': user_2d[:, 0],
    'y': user_2d[:, 1],
    'cluster': sample_labels
})

plt.figure(figsize=(10, 6))
sns.scatterplot(data=df_plot, x='x', y='y', hue='cluster', palette='tab10', s=50)
plt.title("User Clusters Visualized via PCA")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.legend(title='Cluster')
plt.savefig(f"{plots_location}\\plots\\user_clusters_k_{k}.png")
plt.show()
```



Precision@4 and Recall@4

```
[251]: # assign cluster to users
clusters_df = pd.DataFrame({
    'user_id' : [index_to_user[i] for i in range(len(cluster_labels))],
    'cluster': cluster_labels
})

clusters_train_df = train_df.merge(clusters_df, on="user_id")
```

```
[252]: clusters_train_df.head()
```

```
[252]:
```

	user_id	anime_id	rating	cluster
0	287522	522	8.0	2
1	447249	2033	7.0	1
2	446384	31318	10.0	3
3	400930	3712	8.0	0
4	153921	37520	8.0	0

```
[253]: # calculate mean rating for each anime across a cluster
cluster_mean_ratings_df = clusters_train_df.drop(columns=["user_id"]) \
    .groupby(["cluster", "anime_id"]).mean().reset_index()
```

```
[254]: precisions, recalls = [], []
for user_id in sample_index_to_user.values():
    p, r = precision_recall_at_k(user_id, k=5)
    if p is not None:
        precisions.append(p)
        recalls.append(r)

mean_precision = np.mean(precisions)
mean_recall = np.mean(recalls)

precision_scores.append(mean_precision)
recall_scores.append(mean_recall)

print(f"Precision@4: {mean_precision:.6f}")
print(f"Recall@4: {mean_recall:.6f}")
```

Precision@4: 0.006948

Recall@4: 0.000376

Training K = 5

```
[255]: # create cluster
k = 5
kmeans = KMeans(n_clusters=k, random_state=42, max_iter=300)
```

```
cluster_labels = kmeans.fit_predict(user_item_sparse)

inertia.append(kmeans.inertia_)
```

Assessing K = 5

```
[256]: sample_labels = cluster_labels[sample_indices]
```

Silhouette

```
[257]: sil_score = silhouette_score(sample_matrix, sample_labels, metric='cosine')
print(f"Silhouette Score: {sil_score:.6f}")
silhoutte_scores.append(sil_score)
```

Silhouette Score: -0.029378

Davies-Bouldin Index

```
[258]: db_score = davies_bouldin_score(sample_matrix.toarray(), sample_labels)
print(f"Davies-Bouldin Index: {db_score:.6f}")
davies_bouldin_scores.append(db_score)
```


Davies-Bouldin Index: 4.887560

Visualize clusters

```
[259]: # visualize inspect the clusters
pca = PCA(n_components=2)
user_2d = pca.fit_transform(sample_matrix.toarray())

df_plot = pd.DataFrame({
    'x': user_2d[:, 0],
    'y': user_2d[:, 1],
    'cluster': sample_labels
})

plt.figure(figsize=(10, 6))
sns.scatterplot(data=df_plot, x='x', y='y', hue='cluster', palette='tab10',
               s=50)
plt.title("User Clusters Visualized via PCA")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.legend(title='Cluster')
plt.savefig(f"{plots_location}\\plots\\user_clusters_k_{k}.png")
plt.show()
```



Precision@5 and Recall@5

```
[260]: # assign cluster to users
clusters_df = pd.DataFrame({
    'user_id' : [index_to_user[i] for i in range(len(cluster_labels))],
    'cluster': cluster_labels
})

clusters_train_df = train_df.merge(clusters_df, on="user_id")
```

```
[261]: clusters_train_df.head()
```

```
[261]:
```

	user_id	anime_id	rating	cluster
0	287522	522	8.0	2
1	447249	2033	7.0	4
2	446384	31318	10.0	3
3	400930	3712	8.0	0
4	153921	37520	8.0	3

```
[262]: # calculate mean rating for each anime across a cluster
cluster_mean_ratings_df = clusters_train_df.drop(columns=["user_id"]) \
    .groupby(["cluster", "anime_id"]).mean().reset_index()
```

```
[263]: precisions, recalls = [], []
for user_id in sample_index_to_user.values():
```

```

    p, r = precision_recall_at_k(user_id, k=5)
    if p is not None:
        precisions.append(p)
        recalls.append(r)

mean_precision = np.mean(precisions)
mean_recall = np.mean(recalls)

precision_scores.append(mean_precision)
recall_scores.append(mean_recall)

print(f"Precision@5: {mean_precision:.6f}")
print(f"Recall@5: {mean_recall:.6f}")

```

Precision@5: 0.003828
Recall@5: 0.000155

Training K = 6

```

[264]: # create cluster
        k = 6
        kmeans = KMeans(n_clusters=k, random_state=42, max_iter=300)
        cluster_labels = kmeans.fit_predict(user_item_sparse)

        inertia.append(kmeans.inertia_)

```

Assessing K = 6

```

[265]: sample_labels = cluster_labels[sample_indices]

```

Silhouette

```

[266]: sil_score = silhouette_score(sample_matrix, sample_labels, metric='cosine')
        print(f"Silhouette Score: {sil_score:.6f}")
        silhoutte_scores.append(sil_score)

```

Silhouette Score: -0.031468

Davies-Bouldin Index

```

[267]: db_score = davies_bouldin_score(sample_matrix.toarray(), sample_labels)
        print(f"Davies-Bouldin Index: {db_score:.6f}")
        davies_bouldin_scores.append(db_score)

```

Davies-Bouldin Index: 5.011471

Visualize clusters

```

[268]: # visualize inspect the clusters
        pca = PCA(n_components=2)
        user_2d = pca.fit_transform(sample_matrix.toarray())

        df_plot = pd.DataFrame({


```

```

    'x': user_2d[:, 0],
    'y': user_2d[:, 1],
    'cluster': sample_labels
})

plt.figure(figsize=(10, 6))
sns.scatterplot(data=df_plot, x='x', y='y', hue='cluster', palette='tab10', s=50)
plt.title("User Clusters Visualized via PCA")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.legend(title='Cluster')
plt.savefig(f"{plots_location}\\plots\\user_clusters_k_{k}.png")
plt.show()

```



Precision@6 and Recall@6

```

[269]: # assign cluster to users
clusters_df = pd.DataFrame({
    'user_id' : [index_to_user[i] for i in range(len(cluster_labels))],
    'cluster': cluster_labels
})

clusters_train_df = train_df.merge(clusters_df, on="user_id")

```



```
[270]: clusters_train_df.head()
```

```
[270]:
```

	user_id	anime_id	rating	cluster
0	287522	522	8.0	2
1	447249	2033	7.0	4
2	446384	31318	10.0	5
3	400930	3712	8.0	0
4	153921	37520	8.0	5

```
[271]: # calculate mean rating for each anime across a cluster
cluster_mean_ratings_df = clusters_train_df.drop(columns=["user_id"]) \
    .groupby(["cluster", "anime_id"]).mean().reset_index()
```

```
[272]: precisions, recalls = [], []
for user_id in sample_index_to_user.values():
    p, r = precision_recall_at_k(user_id, k=5)
    if p is not None:
        precisions.append(p)
        recalls.append(r)
```

```
mean_precision = np.mean(precisions)
mean_recall = np.mean(recalls)
```

```
precision_scores.append(mean_precision)
recall_scores.append(mean_recall)
```

```
print(f"Precision@6: {mean_precision:.6f}")
print(f"Recall@6: {mean_recall:.6f}")
```

Precision@6: 0.001418

Recall@6: 0.000053

Training K = 7

```
[273]: # create cluster
k = 7
kmeans = KMeans(n_clusters=k, random_state=42, max_iter=300)
cluster_labels = kmeans.fit_predict(user_item_sparse)

inertia.append(kmeans.inertia_)
```

Assessing K = 7

```
[274]: sample_labels = cluster_labels[sample_indices]
```

Silhouette

```
[275]: sil_score = silhouette_score(sample_matrix, sample_labels, metric='cosine')
print(f"Silhouette Score: {sil_score:.6f}")
silhoutte_scores.append(sil_score)
```

Silhouette Score: -0.032445

Davies-Bouldin Index

```
[276]: db_score = davies_bouldin_score(sample_matrix.toarray(), sample_labels)
print(f"Davies-Bouldin Index: {db_score:.6f}")
davies_bouldin_scores.append(db_score)
```


Davies-Bouldin Index: 4.883651

Visualize clusters

```
[277]: # visualize inspect the clusters
pca = PCA(n_components=2)
user_2d = pca.fit_transform(sample_matrix.toarray())

df_plot = pd.DataFrame({
    'x': user_2d[:, 0],
    'y': user_2d[:, 1],
    'cluster': sample_labels
})

plt.figure(figsize=(10, 6))
sns.scatterplot(data=df_plot, x='x', y='y', hue='cluster', palette='tab10', s=50)
plt.title("User Clusters Visualized via PCA")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.legend(title='Cluster')
plt.savefig(f"{plots_location}\\plots\\user_clusters_k_{k}.png")
plt.show()
```



Precision@7 and Recall@7

```
[278]: # assign cluster to users
clusters_df = pd.DataFrame({
    'user_id' : [index_to_user[i] for i in range(len(cluster_labels))],
    'cluster': cluster_labels
})

clusters_train_df = train_df.merge(clusters_df, on="user_id")
```

```
[279]: clusters_train_df.head()
```

```
[279]:
```

	user_id	anime_id	rating	cluster
0	287522	522	8.0	2
1	447249	2033	7.0	4
2	446384	31318	10.0	5
3	400930	3712	8.0	0
4	153921	37520	8.0	5

```
[280]: # calculate mean rating for each anime across a cluster
cluster_mean_ratings_df = clusters_train_df.drop(columns=["user_id"]) \
    .groupby(["cluster", "anime_id"]).mean().reset_index()
```

```
[281]: precisions, recalls = [], []
for user_id in sample_index_to_user.values():
```

```

p, r = precision_recall_at_k(user_id, k=5)
if p is not None:
    precisions.append(p)
    recalls.append(r)

mean_precision = np.mean(precisions)
mean_recall = np.mean(recalls)

precision_scores.append(mean_precision)
recall_scores.append(mean_recall)

print(f"Precision@7: {mean_precision:.6f}")
print(f"Recall@7: {mean_recall:.6f}")

```

Precision@7: 0.001276
Recall@7: 0.000047

Training K = 8

```

[282]: # create cluster
k = 8
kmeans = KMeans(n_clusters=k, random_state=42, max_iter=300)
cluster_labels = kmeans.fit_predict(user_item_sparse)

inertia.append(kmeans.inertia_)

```

Assessing K = 8

```

[283]: sample_labels = cluster_labels[sample_indices]

```

Silhouette

```

[284]: sil_score = silhouette_score(sample_matrix, sample_labels, metric='cosine')
print(f"Silhouette Score: {sil_score:.6f}")
silhoutte_scores.append(sil_score)

```

Silhouette Score: -0.032054

Davies-Bouldin Index

```

[285]: db_score = davies_bouldin_score(sample_matrix.toarray(), sample_labels)
print(f"Davies-Bouldin Index: {db_score:.6f}")
davies_bouldin_scores.append(db_score)

```

Davies-Bouldin Index: 5.124099

Visualize clusters

```

[286]: # visualize inspect the clusters
pca = PCA(n_components=2)
user_2d = pca.fit_transform(sample_matrix.toarray())

df_plot = pd.DataFrame({


```

```

    'x': user_2d[:, 0],
    'y': user_2d[:, 1],
    'cluster': sample_labels
})

plt.figure(figsize=(10, 6))
sns.scatterplot(data=df_plot, x='x', y='y', hue='cluster', palette='tab10', s=50)
plt.title("User Clusters Visualized via PCA")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.legend(title='Cluster')
plt.savefig(f"{plots_location}\\plots\\user_clusters_k_{k}.png")
plt.show()

```



Precision@8 and Recall@8

```

[287]: # assign cluster to users
clusters_df = pd.DataFrame({
    'user_id' : [index_to_user[i] for i in range(len(cluster_labels))],
    'cluster': cluster_labels
})

clusters_train_df = train_df.merge(clusters_df, on="user_id")

```

```
[288]: clusters_train_df.head()
```

```
[288]:
```

	user_id	anime_id	rating	cluster
0	287522	522	8.0	2
1	447249	2033	7.0	1
2	446384	31318	10.0	4
3	400930	3712	8.0	0
4	153921	37520	8.0	4

```
[289]: # calculate mean rating for each anime across a cluster
cluster_mean_ratings_df = clusters_train_df.drop(columns=["user_id"]) \
    .groupby(["cluster", "anime_id"]).mean().reset_index()
```

```
[290]: precisions, recalls = [], []
for user_id in sample_index_to_user.values():
    p, r = precision_recall_at_k(user_id, k=5)
    if p is not None:
        precisions.append(p)
        recalls.append(r)
```

```
mean_precision = np.mean(precisions)
mean_recall = np.mean(recalls)
```

```
precision_scores.append(mean_precision)
recall_scores.append(mean_recall)
```

```
print(f"Precision@8: {mean_precision:.6f}")
print(f"Recall@8: {mean_recall:.6f}")
```

Precision@8: 0.001134

Recall@8: 0.000042

Training K = 9

```
[291]: # create cluster
k = 9
kmeans = KMeans(n_clusters=k, random_state=42, max_iter=300)
cluster_labels = kmeans.fit_predict(user_item_sparse)

inertia.append(kmeans.inertia_)
```

Assessing K = 9

```
[292]: sample_labels = cluster_labels[sample_indices]
```

Silhouette

```
[293]: sil_score = silhouette_score(sample_matrix, sample_labels, metric='cosine')
print(f"Silhouette Score: {sil_score:.6f}")
silhoutte_scores.append(sil_score)
```

Silhouette Score: -0.034246

Davies-Bouldin Index

```
[294]: db_score = davies_bouldin_score(sample_matrix.toarray(), sample_labels)
print(f"Davies-Bouldin Index: {db_score:.6f}")
davies_bouldin_scores.append(db_score)
```


Davies-Bouldin Index: 4.943429

Visualize clusters

```
[295]: # visualize inspect the clusters
pca = PCA(n_components=2)
user_2d = pca.fit_transform(sample_matrix.toarray())

df_plot = pd.DataFrame({
    'x': user_2d[:, 0],
    'y': user_2d[:, 1],
    'cluster': sample_labels
})

plt.figure(figsize=(10, 6))
sns.scatterplot(data=df_plot, x='x', y='y', hue='cluster', palette='tab10',
               s=50)
plt.title("User Clusters Visualized via PCA")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.legend(title='Cluster')
plt.savefig(f"{plots_location}\\plots\\user_clusters_k_{k}.png")
plt.show()
```



Precision@9 and Recall@9

```
[296]: # assign cluster to users
clusters_df = pd.DataFrame({
    'user_id' : [index_to_user[i] for i in range(len(cluster_labels))],
    'cluster': cluster_labels
})

clusters_train_df = train_df.merge(clusters_df, on="user_id")
```

```
[297]: clusters_train_df.head()
```

```
[297]:
```

	user_id	anime_id	rating	cluster
0	287522	522	8.0	2
1	447249	2033	7.0	8
2	446384	31318	10.0	4
3	400930	3712	8.0	0
4	153921	37520	8.0	6

```
[298]: # calculate mean rating for each anime across a cluster
cluster_mean_ratings_df = clusters_train_df.drop(columns=["user_id"]) \
    .groupby(["cluster", "anime_id"]).mean().reset_index()
```

```
[299]: precisions, recalls = [], []
for user_id in sample_index_to_user.values():
```



```

    p, r = precision_recall_at_k(user_id, k=5)
    if p is not None:
        precisions.append(p)
        recalls.append(r)

mean_precision = np.mean(precisions)
mean_recall = np.mean(recalls)

precision_scores.append(mean_precision)
recall_scores.append(mean_recall)

print(f"Precision@9: {mean_precision:.6f}")
print(f"Recall@9: {mean_recall:.6f}")

```

Precision@9: 0.000496
Recall@9: 0.000016

Training K = 10

```

[300]: # create cluster
k = 10
kmeans = KMeans(n_clusters=k, random_state=42, max_iter=300)
cluster_labels = kmeans.fit_predict(user_item_sparse)

inertia.append(kmeans.inertia_)

```

Assessing K = 10

```

[301]: sample_labels = cluster_labels[sample_indices]

```

Silhouette

```

[302]: sil_score = silhouette_score(sample_matrix, sample_labels, metric='cosine')
print(f"Silhouette Score: {sil_score:.6f}")
silhoutte_scores.append(sil_score)

```

Silhouette Score: -0.033874

Davies-Bouldin Index

```

[303]: db_score = davies_bouldin_score(sample_matrix.toarray(), sample_labels)
print(f"Davies-Bouldin Index: {db_score:.6f}")
davies_bouldin_scores.append(db_score)

```

Davies-Bouldin Index: 5.136440

Visualize clusters

```

[304]: # visualize inspect the clusters
pca = PCA(n_components=2)
user_2d = pca.fit_transform(sample_matrix.toarray())

df_plot = pd.DataFrame({


```

```

    'x': user_2d[:, 0],
    'y': user_2d[:, 1],
    'cluster': sample_labels
})

plt.figure(figsize=(10, 6))
sns.scatterplot(data=df_plot, x='x', y='y', hue='cluster', palette='tab10',
               s=50)
plt.title("User Clusters Visualized via PCA")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.legend(title='Cluster')
plt.savefig(f"{plots_location}\\plots\\user_clusters_k_{k}.png")
plt.show()

```



Precision@10 and Recall@10

```

[305]: # assign cluster to users
clusters_df = pd.DataFrame({
    'user_id' : [index_to_user[i] for i in range(len(cluster_labels))],
    'cluster': cluster_labels
})

clusters_train_df = train_df.merge(clusters_df, on="user_id")

```

```
[306]: clusters_train_df.head()
```

```
[306]:
```

	user_id	anime_id	rating	cluster
0	287522	522	8.0	2
1	447249	2033	7.0	2
2	446384	31318	10.0	6
3	400930	3712	8.0	0
4	153921	37520	8.0	6

```
[307]: # calculate mean rating for each anime across a cluster
cluster_mean_ratings_df = clusters_train_df.drop(columns=["user_id"]) \
    .groupby(["cluster", "anime_id"]).mean().reset_index()
```

```
[308]: precisions, recalls = [], []
for user_id in sample_index_to_user.values():
    p, r = precision_recall_at_k(user_id, k=5)
    if p is not None:
        precisions.append(p)
        recalls.append(r)
```

```
mean_precision = np.mean(precisions)
mean_recall = np.mean(recalls)
```

```
precision_scores.append(mean_precision)
recall_scores.append(mean_recall)
```

```
print(f"Precision@10: {mean_precision:.6f}")
print(f"Recall@10: {mean_recall:.6f}")
```

Precision@10: 0.000425

Recall@10: 0.000011

Training K = 11

```
[309]: # create cluster
k = 11
kmeans = KMeans(n_clusters=k, random_state=42, max_iter=300)
cluster_labels = kmeans.fit_predict(user_item_sparse)

inertia.append(kmeans.inertia_)
```

Assessing K = 11

```
[310]: sample_labels = cluster_labels[sample_indices]
```

Silhouette

```
[311]: sil_score = silhouette_score(sample_matrix, sample_labels, metric='cosine')
print(f"Silhouette Score: {sil_score:.6f}")
silhoutte_scores.append(sil_score)
```

Silhouette Score: -0.033710

Davies-Bouldin Index

```
[312]: db_score = davies_bouldin_score(sample_matrix.toarray(), sample_labels)
print(f"Davies-Bouldin Index: {db_score:.6f}")
davies_bouldin_scores.append(db_score)
```


Davies-Bouldin Index: 5.098260

Visualize clusters

```
[313]: # visualize inspect the clusters
pca = PCA(n_components=2)
user_2d = pca.fit_transform(sample_matrix.toarray())

df_plot = pd.DataFrame({
    'x': user_2d[:, 0],
    'y': user_2d[:, 1],
    'cluster': sample_labels
})

plt.figure(figsize=(10, 6))
sns.scatterplot(data=df_plot, x='x', y='y', hue='cluster', palette='tab10',
               s=50)
plt.title("User Clusters Visualized via PCA")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.legend(title='Cluster')
plt.savefig(f"{plots_location}\\plots\\user_clusters_k_{k}.png")
plt.show()
```



Precision@11 and Recall@11

```
[314]: # assign cluster to users
clusters_df = pd.DataFrame({
    'user_id' : [index_to_user[i] for i in range(len(cluster_labels))],
    'cluster': cluster_labels
})

clusters_train_df = train_df.merge(clusters_df, on="user_id")
```

```
[315]: clusters_train_df.head()
```

```
[315]:
```

	user_id	anime_id	rating	cluster
0	287522	522	8.0	2
1	447249	2033	7.0	1
2	446384	31318	10.0	6
3	400930	3712	8.0	0
4	153921	37520	8.0	6

```
[316]: # calculate mean rating for each anime across a cluster
cluster_mean_ratings_df = clusters_train_df.drop(columns=["user_id"]) \
    .groupby(["cluster", "anime_id"]).mean().reset_index()
```

```
[317]: precisions, recalls = [], []
for user_id in sample_index_to_user.values():
```

```

    p, r = precision_recall_at_k(user_id, k=5)
    if p is not None:
        precisions.append(p)
        recalls.append(r)

mean_precision = np.mean(precisions)
mean_recall = np.mean(recalls)

precision_scores.append(mean_precision)
recall_scores.append(mean_recall)

print(f"Precision@11: {mean_precision:.6f}")
print(f"Recall@11: {mean_recall:.6f}")

```

```

Precision@11: 0.000000
Recall@11: 0.000000

```

Training K = 12

```

[318]: # create cluster
k = 12
kmeans = KMeans(n_clusters=k, random_state=42, max_iter=300)
cluster_labels = kmeans.fit_predict(user_item_sparse)

inertia.append(kmeans.inertia_)

```

Assessing K = 12

```

[319]: sample_labels = cluster_labels[sample_indices]

```

Silhouette

```

[320]: sil_score = silhouette_score(sample_matrix, sample_labels, metric='cosine')
print(f"Silhouette Score: {sil_score:.6f}")
silhoutte_scores.append(sil_score)

```

```

Silhouette Score: -0.034318

```

Davies-Bouldin Index

```

[321]: db_score = davies_bouldin_score(sample_matrix.toarray(), sample_labels)
print(f"Davies-Bouldin Index: {db_score:.6f}")
davies_bouldin_scores.append(db_score)

```

```

Davies-Bouldin Index: 4.960168

```

Visualize clusters

```

[322]: # visualize inspect the clusters
pca = PCA(n_components=2)
user_2d = pca.fit_transform(sample_matrix.toarray())

df_plot = pd.DataFrame({


```

```

    'x': user_2d[:, 0],
    'y': user_2d[:, 1],
    'cluster': sample_labels
})

plt.figure(figsize=(10, 6))
sns.scatterplot(data=df_plot, x='x', y='y', hue='cluster', palette='tab10',
               s=50)
plt.title("User Clusters Visualized via PCA")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.legend(title='Cluster')
plt.savefig(f"{plots_location}\\plots\\user_clusters_k_{k}.png")
plt.show()

```



Precision@12 and Recall@12

```

[323]: # assign cluster to users
clusters_df = pd.DataFrame({
    'user_id' : [index_to_user[i] for i in range(len(cluster_labels))],
    'cluster': cluster_labels
})

clusters_train_df = train_df.merge(clusters_df, on="user_id")

```

```
[324]: clusters_train_df.head()
```

```
[324]:
```

	user_id	anime_id	rating	cluster
0	287522	522	8.0	7
1	447249	2033	7.0	11
2	446384	31318	10.0	6
3	400930	3712	8.0	0
4	153921	37520	8.0	6

```
[325]: # calculate mean rating for each anime across a cluster
cluster_mean_ratings_df = clusters_train_df.drop(columns=["user_id"]) \
    .groupby(["cluster", "anime_id"]).mean().reset_index()
```

```
[326]: precisions, recalls = [], []
for user_id in sample_index_to_user.values():
    p, r = precision_recall_at_k(user_id, k=5)
    if p is not None:
        precisions.append(p)
        recalls.append(r)

mean_precision = np.mean(precisions)
mean_recall = np.mean(recalls)

precision_scores.append(mean_precision)
recall_scores.append(mean_recall)

print(f"Precision@12: {mean_precision:.6f}")
print(f"Recall@12: {mean_recall:.6f}")
```

Precision@12: 0.001347

Recall@12: 0.000090

Assessing Inertia, Silhouette and Davies-Bouldin Scores

```
[ ]: k_values = range(2, 13)

plt.figure(figsize=(28, 12))
plt.subplot(3, 1, 1)
plt.plot(k_values, inertia, marker='o')
plt.title('KMeans Inertia for Different Values of k')
plt.xlabel('Number of Clusters (k)')
plt.ylabel('Inertia')
plt.xticks(k_values)
plt.grid(True)

plt.subplot(3, 1, 2)
plt.plot(k_values, silhouette_scores, marker='o', color='orange')
plt.title('Silhouette Scores for Different Values of k')
```




```

plt.xlabel('Number of Clusters (k)')
plt.ylabel('Silhouette Score')
plt.xticks(k_values)
plt.grid(True)

plt.subplot(3, 1, 3)
plt.plot(k_values, davies_bouldin_scores, marker='o', color='pink')
plt.title('Davies-Bouldin Score for Different Values of k')
plt.xlabel('Number of Clusters (k)')
plt.ylabel('Davies-Bouldin Score')
plt.xticks(k_values)
plt.grid(True)

plt.tight_layout()
plt.savefig(f"{plots_location}\\plots\\inertia_silhouette_davies_bouldin.png")
plt.show()

```



Comparing Precision@K and Recall@K

```

[328]: recommendation_metrics_df = pd.DataFrame(columns=["K", "Precision", "Recall"])
recommendation_metrics_df["K"] = [i for i in range(2, 13)]
recommendation_metrics_df["Precision"] = precision_scores
recommendation_metrics_df["Recall"] = recall_scores

```

```

[330]: recommendation_metrics_df

```

```

[330]:
   K  Precision  Recall
0  2    0.017795  0.001476
1  3    0.007657  0.000430
2  4    0.006948  0.000376
3  5    0.003828  0.000155

```

4	6	0.001418	0.000053
5	7	0.001276	0.000047
6	8	0.001134	0.000042
7	9	0.000496	0.000016
8	10	0.000425	0.000011
9	11	0.000000	0.000000
10	12	0.001347	0.000090

```
[332]: recommendation_metrics_df.  
        ↪to_csv(f"{plots_location}\\metrics\\precision_and_recall.csv", index=False)
```

k-means @ 4 is chosen for the model as it provides the best clusters based on visual inspection as well as the inertia, silhouette and davies-bouldin scores

```
[ ]: import joblib  
      # create cluster  
      k = 4  
      kmeans = KMeans(n_clusters=k, random_state=42, max_iter=300)  
      cluster_labels = kmeans.fit_predict(user_item_sparse)  
  
      clusters_df = pd.DataFrame({  
          'user_id' : [index_to_user[i] for i in range(len(cluster_labels))],  
          'cluster': cluster_labels  
      })  
  
      clusters_train_df = train_df.merge(clusters_df, on="user_id")  
  
      model_bundle = {  
          'kmeans': kmeans,  
          'df': clusters_train_df,  
          'id_to_user': index_to_user,  
          'id_to_anime': index_to_anime  
      }  
  
      # Store to file  
      joblib.dump(model_bundle, "E:\\\\applied data science_\\  
          ↪capstone\\clusters\\user_clustering_model.joblib")  
      #
```